

NIRANJAN VSKS Senior Agentic AI Engineer | Forward Deployed Engineering

✉ niranjan.vsk@gmail.com ☎ +91 7893312735 📍 Hyderabad, India 🔗 niranjanvsk.xyz 🔗 linkedin.com/in/niranjanvsk
📁 github.com/niranjan-vsk

Summary

Senior **Agentic AI Engineer** with **7 years** of expertise. I build **production agentic systems** that run inside the **customer's own cloud**, under their **identity**, their constraints and their **data boundaries**. I own the layer that keeps them **reliable & efficient at scale: evaluation, latency and token cost**. Blank page to production. **ML and statistics underneath**. The engineer who **architects, builds and deploys** from scratch.

Experience

Senior Engineer - GenAI & Agentic AI, Coforge

03/2024 – Present | Hyderabad

- Architected and shipped an **AI-Infused Agentic Quality Engineering Platform** end to end; took it from feasibility framing to production with **17 enterprise QA teams** at peak adoption, and deployed it into customer-managed **AWS, Azure, and GCP** tenants with architecture, compute sizing, and identity (**Azure AD, IAM, SSO**) adapted per cloud, running on **Azure OpenAI and Azure AI Foundry** all on one codebase.
- Cut production hallucination from **~15% to under 5%** with **GraphRAG on Neo4j** over **Jira** and Azure DevOps: **entity normalization, query transformation, re-ranking and context compression**, so generation is scoped against linked stories, epics and **Acceptance criteria** rather than flat text.
- Owned **cost, latency and output quality** as one problem: token and cost engineering (**context caching, prompt-payload reduction, tiered model routing**) underneath a **live evaluation and FinOps layer** with output quality KPIs I defined from scratch (**Acceptance Criteria Coverage, Test Design Coverage, RAGAS**) plus **per-tenant cost telemetry** that stakeholders read directly.
- Eliminated **85-90% of manual QA effort** with agentic test generation, execution, and **CI-integrated automation**; built a custom **Playwright MCP server** that turns natural language into production-grade Playwright and Cypress scripts and cut test authoring time **~75%**.
- Secured **multi-tenant isolation** across clouds with **tenant-aware RBAC**: module-level **CRUD-X** permissions and role templates that let new enterprise customers onboard without a rebuilt permission model.
- The practice's **AI technical authority: L1 technical interviewer** for **AI/ML, GenAI & Agentic AI** hiring across **50+ interviews**, setting the technical bar, and **mentoring** engineers on **agent design, retrieval architecture and evaluation practice**.

Data Scientist - GenAI, ML & Conversational AI, Hewlett Packard Enterprise

11/2022 – 12/2023 | Bangalore

- Architected the migration of an enterprise conversational system from **rule-based FAQ to LLM-assisted RAG** over ~1,500-1,700 internal documents, cutting repeat user escalations **~40%**.
- Designed the retrieval pipeline end to end: chunking with overlap calibration, **hybrid search combining semantic and keyword signals, BM25** and metadata filtering tuned for enterprise-jargon queries. Added evaluation loops that measured retrieval precision against a **curated held-out evaluation set** and improved response relevance 30-40% in pilot.
- Built the **Power BI** reporting layer over conversation and usage data so stakeholders tracked **adoption, deflection and escalation trends** directly, replacing ad hoc reporting requests with a self-serve view.

Associate Software Engineer - Machine Learning & AI, Mphasis

07/2019 – 10/2022 | Bangalore

- Built **fraud-detection and risk-scoring models** for financial services clients using **tree based ensembles (XGBoost, LightGBM, Random Forest)** with **SQL-based feature extraction** over transactional data, improving **precision ~15%** over legacy rule-based systems and cutting manual review volume.
- Handled the **class imbalance** that makes fraud data misleading on a naive split, and engineered the **behavioural and transactional features** the models ran on from raw client data

Education

Product Management in Generative & Agentic AI, BITS School of Management (BITSoM)

07/2025 – 03/2026 | India

PG Diploma in Machine Learning & Artificial Intelligence, IIIT Bangalore

06/2020 – 03/2022 | Bangalore

Bachelor of Technology in Computer Science, Bharat Institute of Engineering and Technology

08/2015 – 06/2019 | Hyderabad

Skills

Generative and Agentic AI: LLMs | Multi-Agent Orchestration | Agentic Workflows | LangGraph | LangChain | MCP Servers | Tool & Function Calling | Structured Outputs | Prompt & Context Engineering | LLM Observability | Token-Cost Engineering | Model Routing | Evaluation Design & Rubrics | RAGAS | Guardrails | Responsible AI

Retrieval & Data Architecture: Knowledge Graphs | GraphRAG | Neo4j | AWS Neptune | Entity Resolution | Hybrid Search | Semantic Querying | Reranking | Embeddings & Vector Search | Agent Memory Architecture | Agentic RAG Pipelines

Enterprise and Customer Facing: Forward Deployed Engineering | Solutions Architecture | Customer Discovery | Stakeholder Alignment | Production Deployment into Customer Tenants (AWS, Azure, GCP) | Multi-Tenant Architecture | Identity & SSO Architecture | RBAC Design | Enterprise API Integration | Customer-Facing Technical Communication

Software Engineering & Cloud: Python | TypeScript | Node.js | FastAPI | Next.js | React | PostgreSQL | MongoDB | REST & Webhook Architecture | Distributed Systems | Azure OpenAI | Azure AI Foundry | Entra ID | AWS (IAM, S3, Neptune) | GCP | Docker | Kubernetes | CI/CD

Projects

Loop Copilot : AI CRM Copilot for Enterprise Sales, Architect · Engineer · Sole Owner | Production, Fortune 500 pilot

- Built and shipped a **production AI CRM copilot** for enterprise sales reps as sole architect and engineer; took it from customer discovery through V2 with an active beta user inside a **Fortune 500 sales org**, cutting CRM logging time from **4-6 min to ~45 sec (~85%)**.

- Shipped a **voice agent pipeline** that listens to sales-client calls (**STT**), turns them into structured **PRD** and **onboarding** documents, and **auto-detects** any newly mentioned accounts. It resolves them against the **account repository** using **unique ID**, then logs them straight into **D365** or the connected **CRM**, so reps never have to manually key in account details after a call.
- Architected the full stack (**React 19, async FastAPI, MongoDB, Azure OpenAI, Entra ID via MSAL, Microsoft Graph**, Telegram Bot API, **CI/CD**) and chose **structured event sourcing** over vector RAG for memory, removing hallucinated recall while keeping **cross-session context**; shipped **10+ feature areas** including calendar integration, bulk upload with AI summarization, and **admin analytics**, with a **multi-CRM** expansion shell (**Salesforce, HubSpot, Zoho**) architected and ready.

Agentic Codebase Intelligence System, Multi-Agent Platform for Enterprise & Architectural Level Codebase Auditing 06/2026 – Present

- Architected and shipped a **multi-agent codebase-intelligence platform** end to end that ingests polyglot repositories and builds a **persistent structural model of the system** (entities, call paths, **cross-file dependencies**) that agents reason over, rather than retrieving flat file text. Cut manual review and audit effort **~50%**
- **Built an audit orchestrator coordinating 6 specialized agents** (retrieval, reasoning, verification) with MCP tool integration, compressing time-to-findings from **~2 days of manual review to under 20 minutes** and auto-filing prioritized issues into Jira.

Autonomous Wealth Operating System, Multi-Asset Agentic Application | Architect. Engineer. Owner | Ongoing

- Architected and built an **autonomous multi-asset wealth platform** (stocks, mutual funds, ETFs, bonds, NPS) around a **21-agent analyst council**: isolated agent opinions, a **devil's advocate stage**, aggregation weighted by each agent's **historical accuracy**, and a **code-enforced veto gate** so no money moves on LLM output alone.
- Engineered a **five-ring risk system** (position, portfolio, drawdown, budget, and connection-level **broker locks**) plus a **frozen-threshold evolution gauntlet**: any strategy change must beat the incumbent across **fixed statistical gates** before promotion, which rules out **silent strategy drift**.
- Delivered it as a **multi-surface platform**: a shared **framework-free core** behind web and mobile clients, a **Telegram copilot** for conversational instructions and live trade, investment, and report updates, and **user-selectable LLM models per agent role**.

Voice-First Financial AI Copilot, Capstone, BITSOM : Architected, Built & Piloted

- Owned architecture and delivery of a **voice-first financial learning and advisory copilot** for India's gig workforce, from customer research through pilot: a **~30-screen MVP** with **10-language architecture**, shipped on both web and mobile.
- Designed a **multi-layer compliance architecture** with **SEBI advisory principles and DPDP requirements** as first-class design constraints: RAG over a regulatory-safe knowledge base, system-prompt guardrails, server-side governance, and consent-first data handling.
- Engineered a **central memory layer with RAG** so user context, goals, and prior conversations persist across sessions and across web and mobile from a single source of truth.

Certifications

Claude Certified Architect	Agile Practices Certificate
Anthropic : In-Progress	SAFe